

# Operating System

*Dr. GuoJun LIU*

Harbin Institute of Technology

<http://guojunos.hit.edu.cn>

# Chapter 04

## *Threads*

线程

# Learning Objectives

---

- Understand the **distinction** between **process** and **thread**
- Describe the **basic design issues** for threads
- Explain the **difference** between **user-level** threads and **kernel-level** threads

# Outline

---

## ■ Processes and Threads

- Multithreading
- Thread Functionality

## ■ Types of Threads

- User-Level and Kernel-Level Threads
- Other Arrangements



The basic idea is that the several components in any complex system will perform **particular subfunctions** that contribute to the **overall function**.

-- *THE SCIENCES OF THE ARTIFICIAL*,  
*Herbert Simon*

Turing Award	1975
Nobel Prize in Economics	1978
National Medal of Science	1986
von Neumann Theory Prize	1988

# Processes and Threads

---

## ■ Resource Ownership

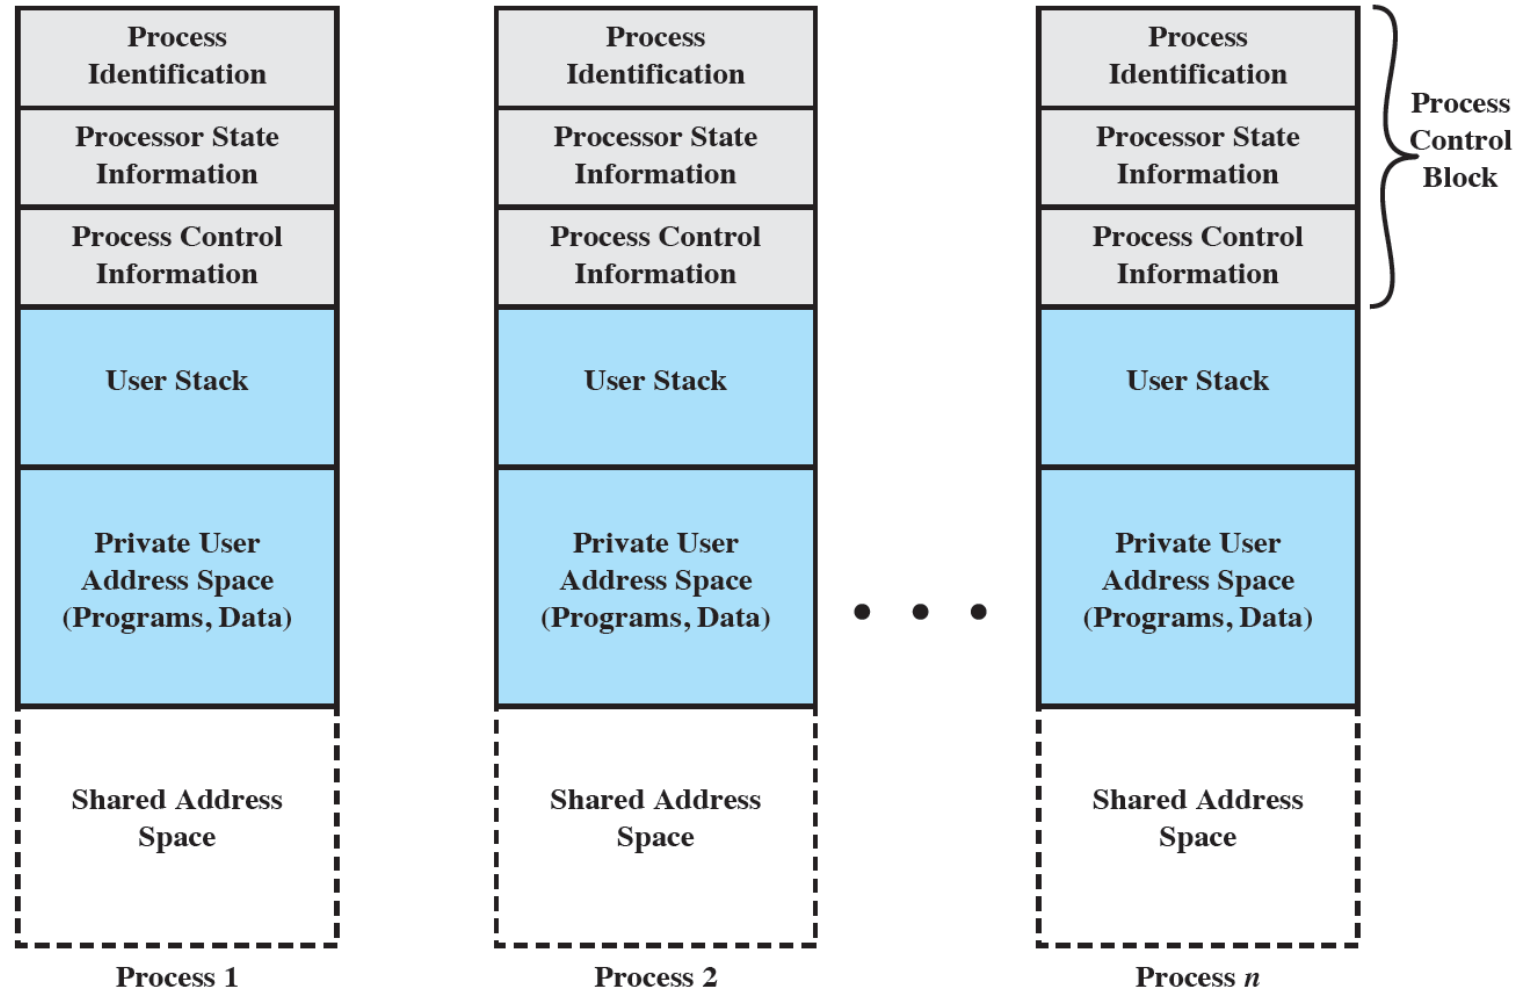
- a **virtual address space** to hold the **process image**
- may be allocated **control** or **ownership** of resources
  - main memory
  - I/O channels and devices
  - files
- OS performs a **protection function** to prevent unwanted interference between processes with respect to resources

## ■ Scheduling/Execution

- Follows an execution path that may be **interleaved** with other processes
- an **execution state**
  - Running, Ready, etc.
- a dispatching **priority**
- is **the entity** that is scheduled and dispatched by the OS



# Structure of Process Images



User Processes in **Virtual Memory**

# Processes and Threads

---

## ■ Thread or lightweight process LWT

- The unit of **dispatching** is referred to as a thread or lightweight process

## ■ Process or Task


- The unit of **resource ownership** is referred to as a process or task

## ■ Multithreading


- The ability of an OS to support multiple, concurrent paths of execution within a single process





# Threads and Processes





one process  
one thread




one process  
multiple threads



multiple processes  
one thread per process



multiple processes  
multiple threads per process



# Processes


---

## ■ The **unit** of resource allocation

- A virtual address space that holds the process image

## ■ The **unit** of protection

- processors
- other processes
  - for interprocess communication
- files
- I/O resources
  - devices and channels



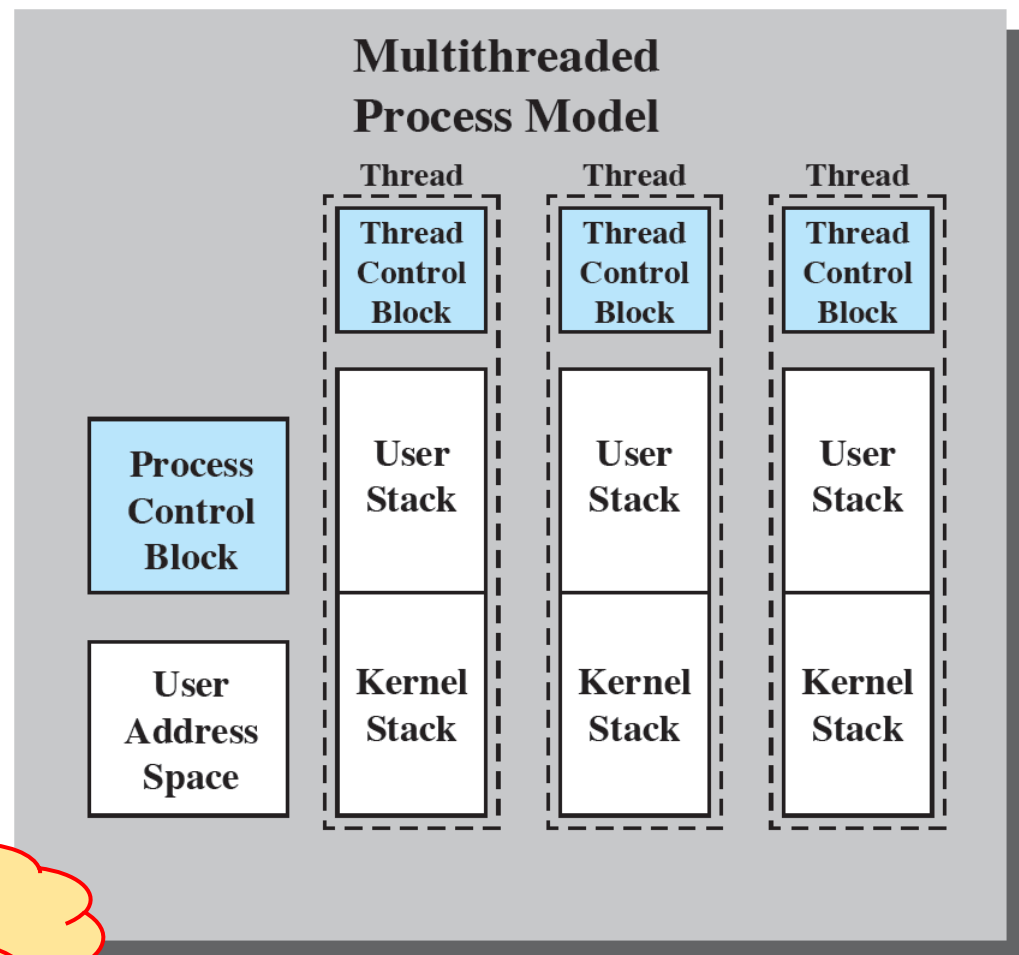
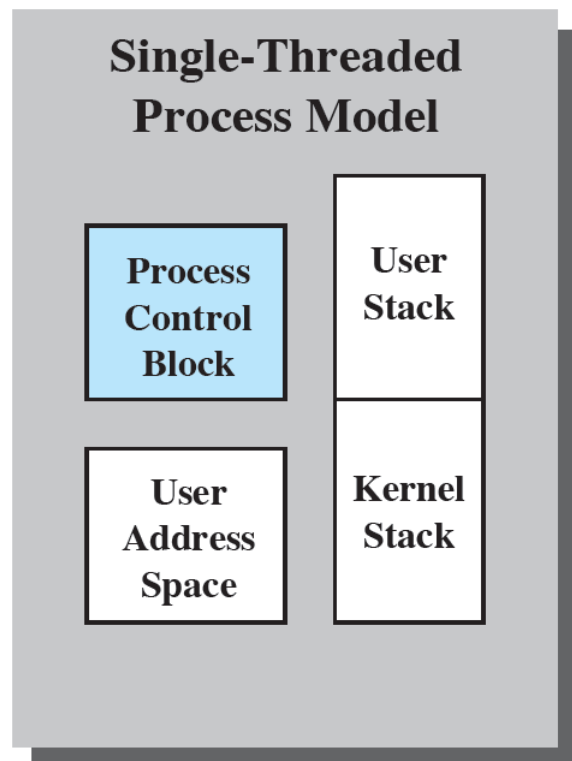
# One or More Threads in a Process

---

## ■ Each thread has

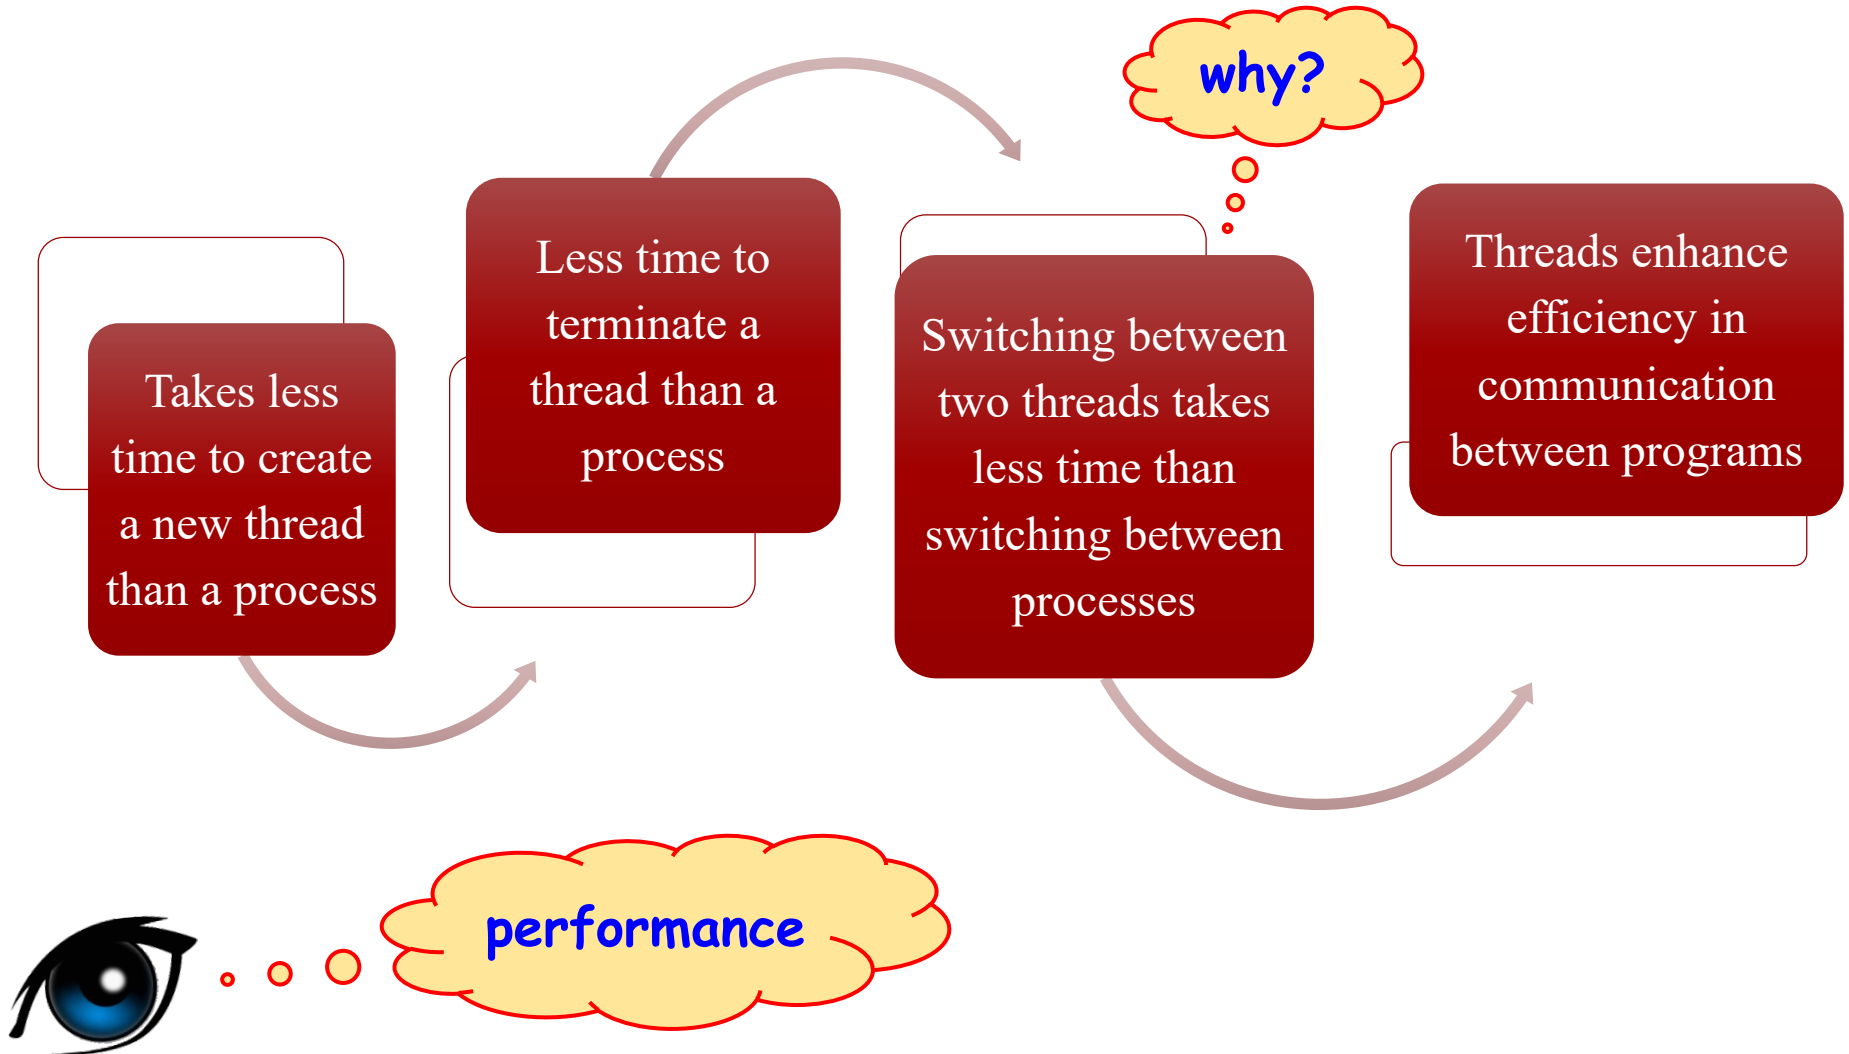
- an execution **state**
  - Running, Ready, etc.
- **saved thread context** when not running
  - one way to view a thread is **as an independent program counter** operating **within a process**
- an execution **stack**
- some **per-thread static storage** for local variables
- access to the **memory and resources of its process**
  - all threads of a process **share** this

# Threads vs. Processes



process  
management

# Key Benefits of Threads



# Thread Use in a Single-User System

---

## ■ Foreground and background work

- a spreadsheet program

## ■ Asynchronous processing

- as a protection against power failure

## ■ Speed of execution

- A multithreaded process can compute one batch of data while reading the next batch from a device

## ■ Modular program structure

- Programs that involve a variety of activities or a variety of sources and destinations of input and output

# Thread Functionality

---

- **Scheduling and dispatching is done on a thread basis**
- **Most of the state information dealing with execution is maintained in thread-level data structures**
- **Several actions that affect all of the threads in a process**
  - the OS must manage at the process level
    - **suspending** a process involves suspending all threads of the process
    - **termination** of a process terminates all threads within the process

# Thread Execution States

---

## ■ Key states for a thread

- Running
- Ready
- Blocked

## ■ Thread operations associated with a change in thread state


- **Spawn**
- Block
- Unblock
- Finish

## ➤ **Spawn**


- when a new process is spawned, a thread for that process is also spawned
- a thread within a process may spawn another thread within the same process
- The new thread is provided with its own register context and stack space and placed on the ready queue




# Performance Benefits of Threads



RPC Using Single Thread



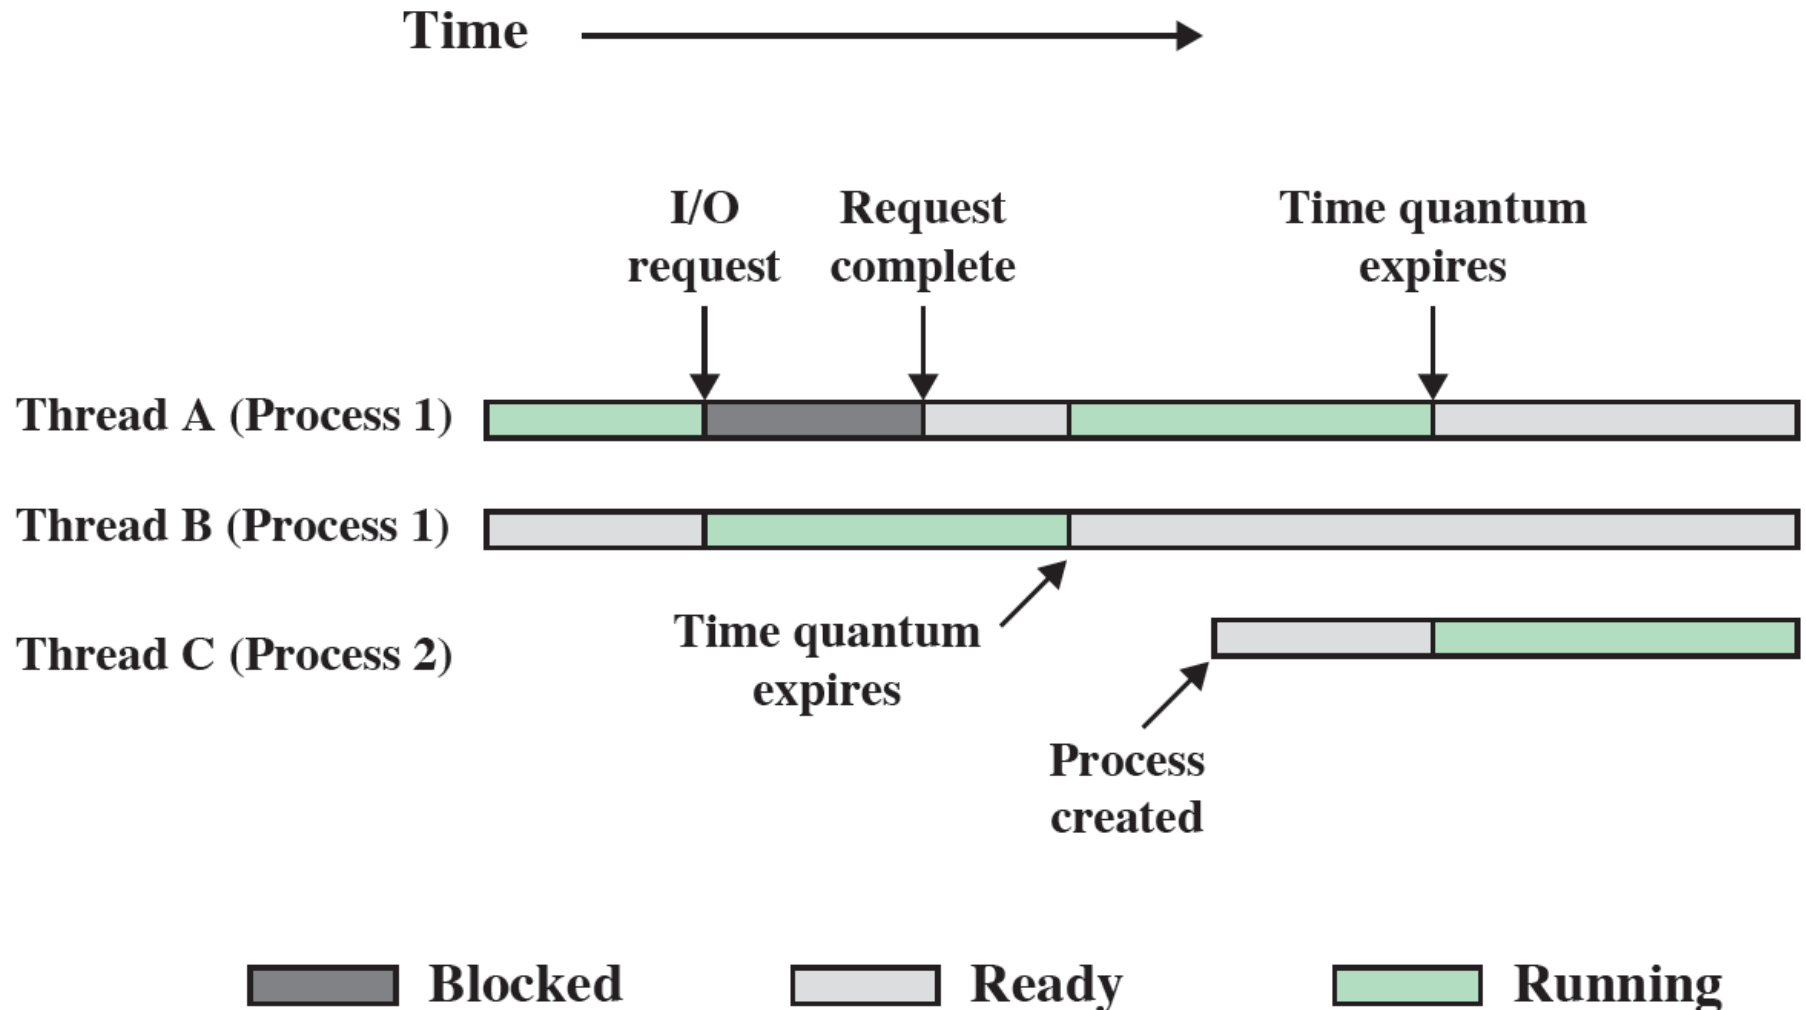
 Blocked, waiting for response to RPC

 Blocked, waiting for processor, which is in use by Thread B

 Running


RPC Using One Thread per Server (on a uniprocessor)

# Multithreading on a Uniprocessor



# Thread Synchronization

- It is necessary to synchronize the activities of the various threads
  - all threads of a process share the same address space and other resources
  - any alteration of a resource by one thread affects the other threads in the same process



if two threads each try to add an element to a doubly linked list at the same time

- one element may be **lost**
- the list may end up **malformed**

# Types of Threads

---

whether the **blocking** of a  
**thread** results in the **blocking**  
of the **entire process** ?

User Level Thread (ULT)

Kernel level Thread (KLT)




# User-Level Threads (ULTs)

- All thread management is done by the **application**

- in user space
- within a single process

- The kernel is not aware of the existence of threads

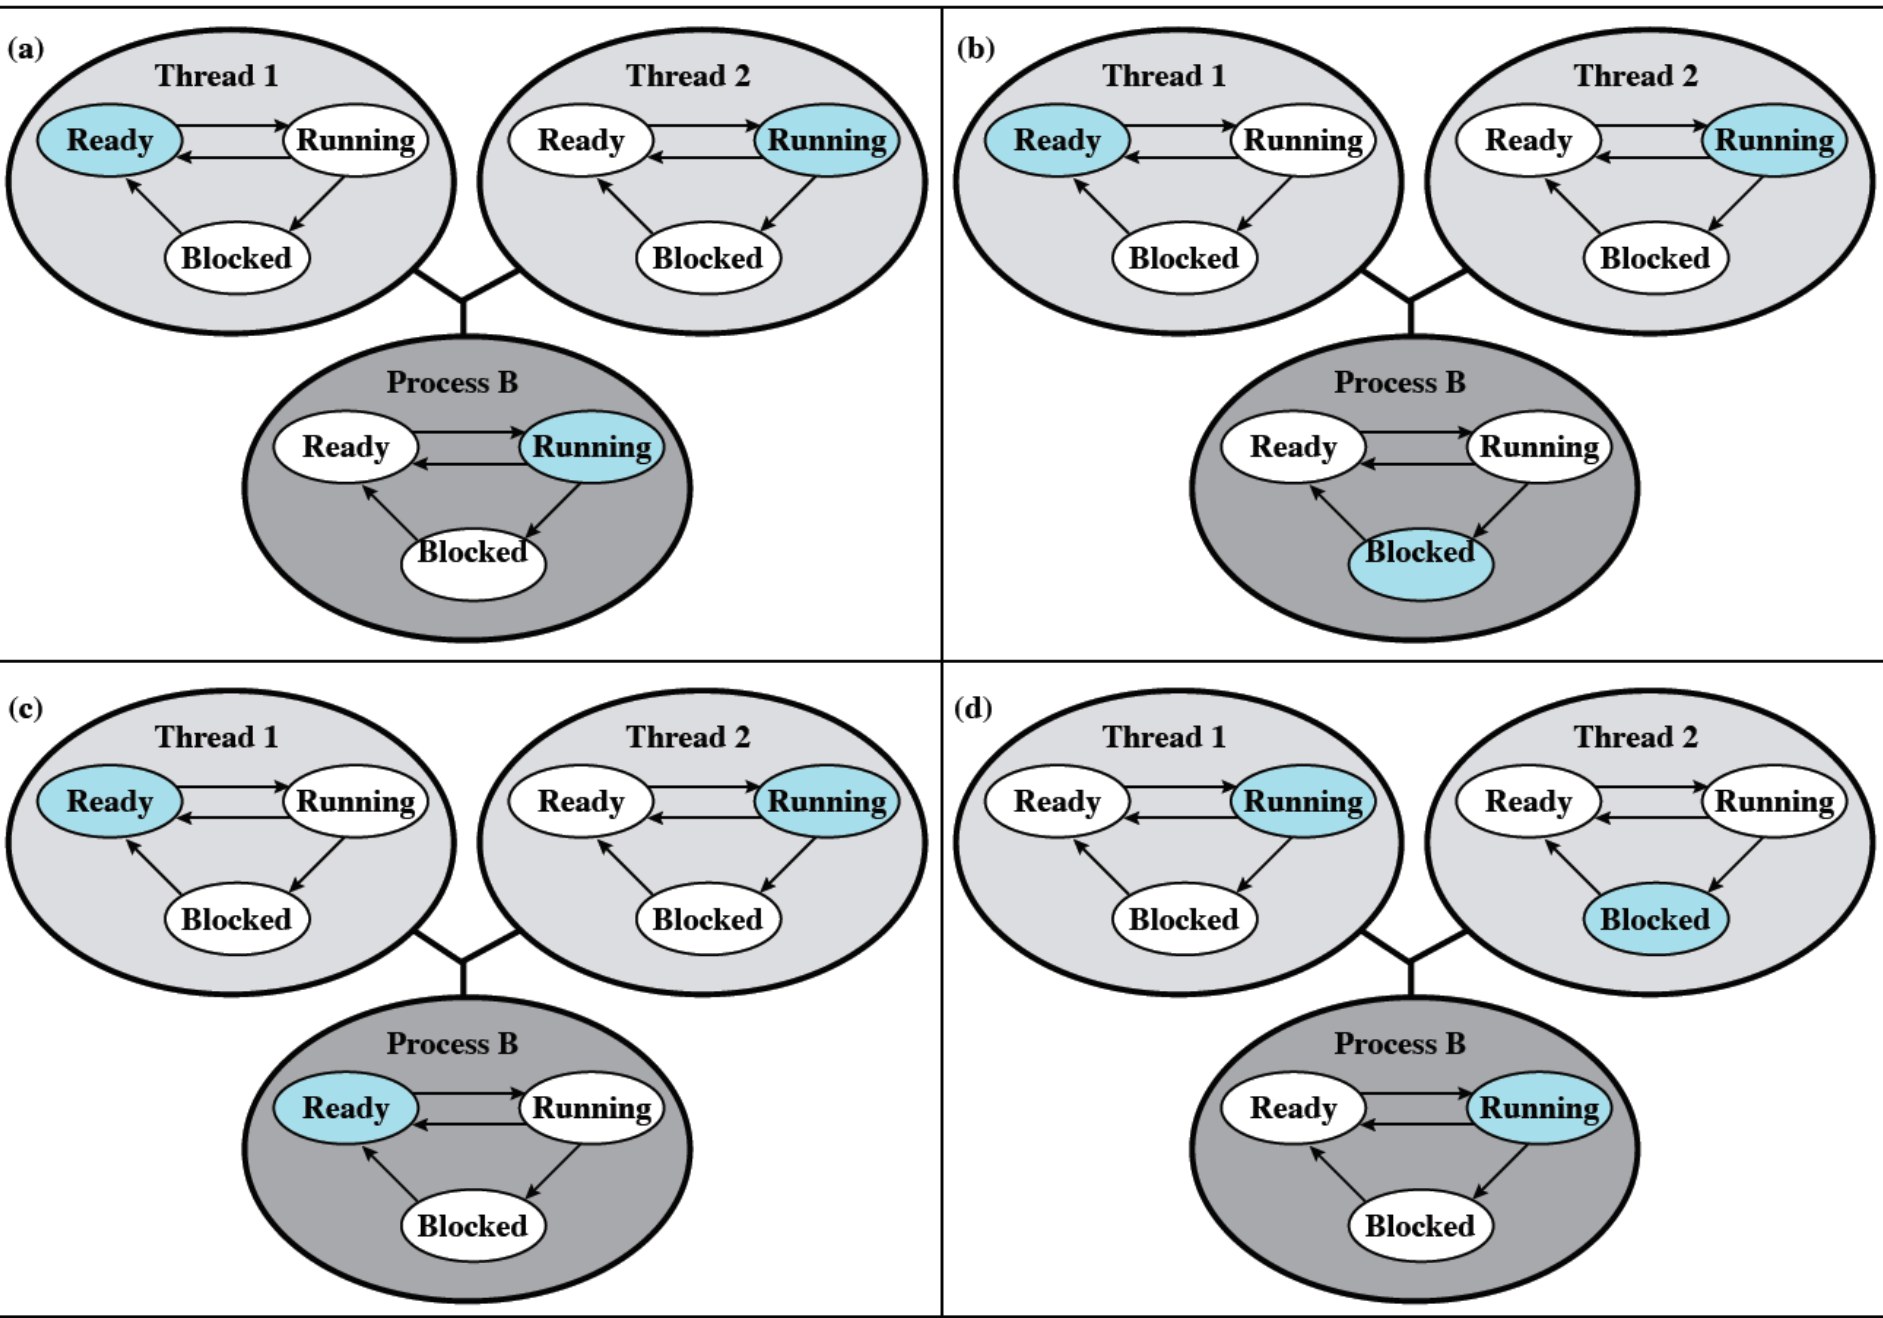


Pure user-level

{ User-level thread

○ Kernel-level thread

○ P Process



Examples of the Relationships Between User-Level Thread States and Process States [Slides-22](#)

# Advantages of ULTs

---

- **Thread switching does not require kernel mode privileges**
  - This **saves the overhead of two mode switches**
    - user to kernel; kernel back to user
- **Scheduling can be application specific**
  - The scheduling algorithm can **be tailored to the application without disturbing** the underlying OS scheduler
- **ULTs can run on any OS**
  - The **threads library** is a set of application-level functions shared by all applications

# Disadvantages of ULTs

---

## ■ Disadvantages

- when a ULT executes a **system call**, **not only** is that thread blocked, **but also all of the threads** within the process are blocked
- In a pure ULT strategy, **a multithreaded application** cannot take advantage of **multiprocessing**

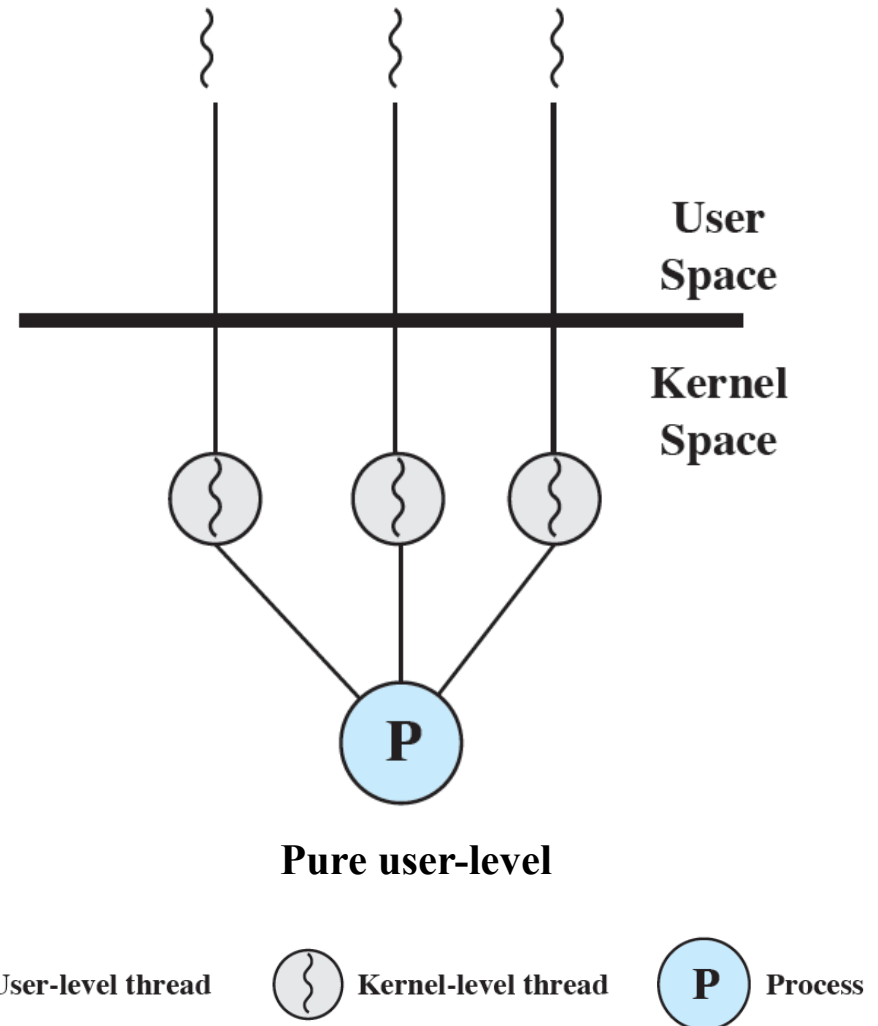
## ■ Overcoming

- Writing an application as **multiple processes** rather than **multiple threads**
  - Each switch becomes a **process switch** rather than a **thread switch**, resulting in much **greater overhead**
- **Jacketing**
  - converts a **blocking** system call into a **non-blocking** system call



# Kernel-Level Threads (KLTs)

- Thread management is **done by the kernel**
  - no thread management is done by the application
  - Windows is an example of this approach



# Advantages of KLTs

---

- The kernel **can simultaneously schedule** multiple threads from the same process **on multiple processors**
- If one thread in a process is blocked, the kernel **can schedule** another thread of the same process
- Kernel routines can be **multithreaded**

# Disadvantage of KLTs

---


- The **transfer** of control from one thread to another within the **same process** requires a mode switch to the kernel

**Table** Thread and Process Operation Latencies ( $\mu\text{s}$ )

Operation	User-Level Threads	Kernel-Level Threads	Processes
Null Fork	34	948	11,300
Signal Wait	37	441	1,840

# Combined Approaches

- **Thread creation** is done in the **user space**
- Bulk of **scheduling** and **synchronization** of threads is by the application
- The multiple ULTs are **mapped onto** some number of KLTs
- The programmer may **adjust the number of KLTs**



Combined



# Relationship Between Threads and Processes

Threads:Processes	Description	Example Systems
<b>1:1</b>	Each thread of execution is a unique process with its own address space and resources.	Traditional UNIX implementations
<b>M:1</b>	A process defines an address space and dynamic resource ownership. Multiple threads may be created and executed within that process.	Windows NT, Solaris, Linux, OS/2, OS/390, MACH
<b>1:M</b>	A thread may migrate from one process environment to another. This allows a thread to be easily moved among distinct systems.	Ra (Clouds), Emerald
<b>M:N</b>	Combines attributes of M:1 and 1:M cases.	TRIX

# Summary

---

## ■ Process

- **resource ownership**

## ■ User-level threads

- created and managed by a threads library that runs in the user space of a process
- a mode switch is not required to switch from one thread to another
- only a single user-level thread within a process can execute at a time
- if one thread blocks, the entire process is blocked

## ■ Thread

- **program execution**

## ■ Kernel-level threads

- threads within a process that are maintained by the kernel
- a mode switch is required to switch from one thread to another
- multiple threads within the same process can execute in parallel on a multiprocessor
- blocking of a thread does not block the entire process