

# Operating System

*Dr. GuoJun LIU*

Harbin Institute of Technology

<http://guojunos.hit.edu.cn>

# Chapter 10

## *I/O Management and Disk Scheduling*

I/O管理和磁盘调度

# Learning Objectives

---

- **Summarize key categories of I/O devices on computers**
- **Discuss the organization of the I/O function**
- **Explain some of the key issues in the design of OS support for I/O**
- **Analyze the performance implications of various I/O buffering alternatives**
- **Understand the performance issues involved in magnetic disk access**

# Outline

---

## ■ I/O Devices

## ■ Organization of the I/O Function

- The Evolution of the I/O Function
- Direct Memory Access

## ■ OS Design Issues

- Design Objectives
- Logical Structure of the I/O Function

## ■ I/O Buffering

- Single Buffer
- Double Buffer
- Circular Buffer
- The Utility of Buffering

## ■ Disk Scheduling

- Disk Performance Parameters
- Disk Scheduling Policies

# Categories of I/O Devices

---

- **External devices that engage in I/O with computer systems can be grouped into three categories:**
  - **Human readable**
    - suitable for communicating with the computer user
    - printers, terminals, video display, keyboard, mouse
  - **Machine readable**
    - suitable for communicating with electronic equipment
    - disk drives, USB keys, sensors, controllers
  - **Communication**
    - suitable for communicating with remote devices
    - modems, digital line drivers

# Differences in I/O Devices

---

## ■ Data Rate

- there may be differences of magnitude between the data transfer rates

## ■ Application

- the use to which a device is put has an influence on the software

## ■ Complexity of Control

- the effect on the operating system is filtered by the complexity of the I/O module that controls the device

## ■ Unit of Transfer

- data may be transferred as a stream of bytes or characters or in larger blocks


## ■ Data Representation

- different data encoding schemes are used by different devices

## ■ Error Conditions

- the nature of errors, the way in which they are reported, their consequences, and the available range of responses differs from one device to another

# Typical I/O Device Data Rates



# Organization of the I/O Function

---

## ■ Programmed I/O

- the processor issues an I/O command on behalf of a process to an I/O module; that process then busy waits for the operation to be completed before proceeding

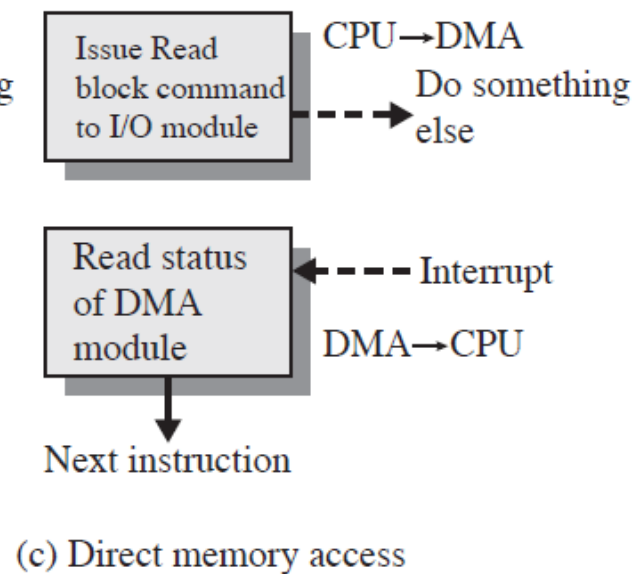
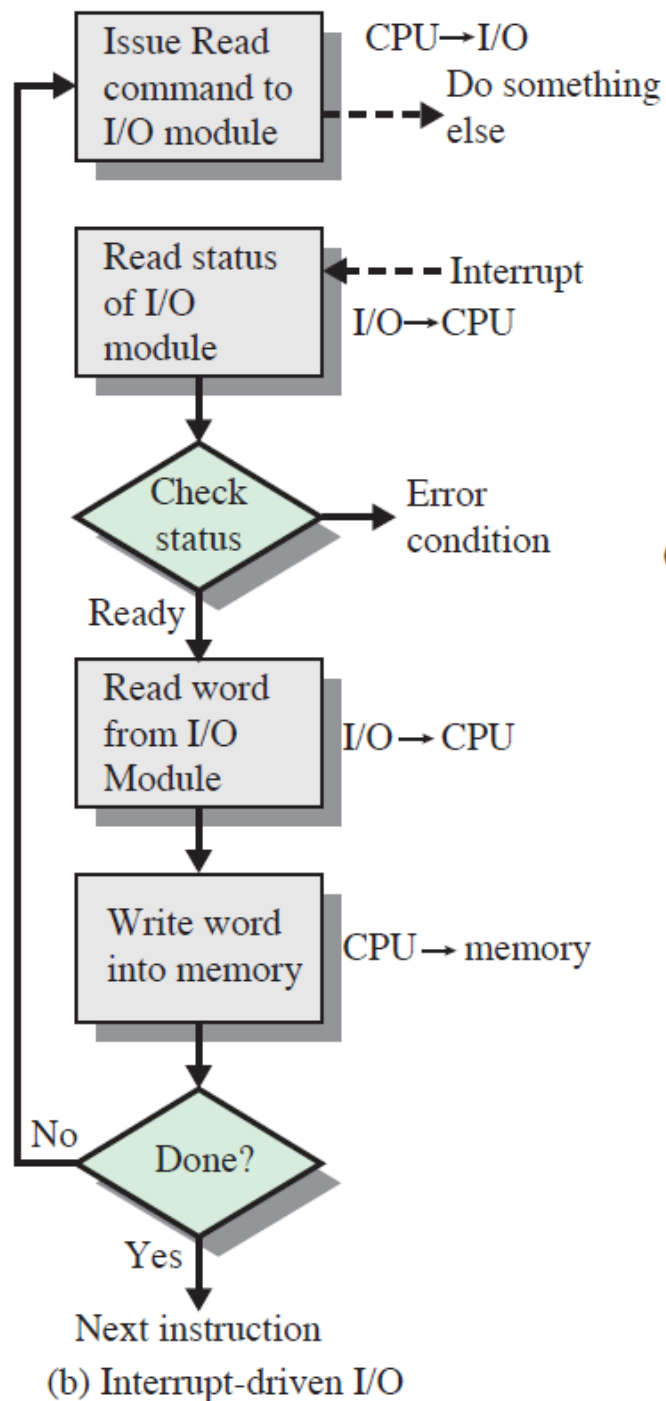
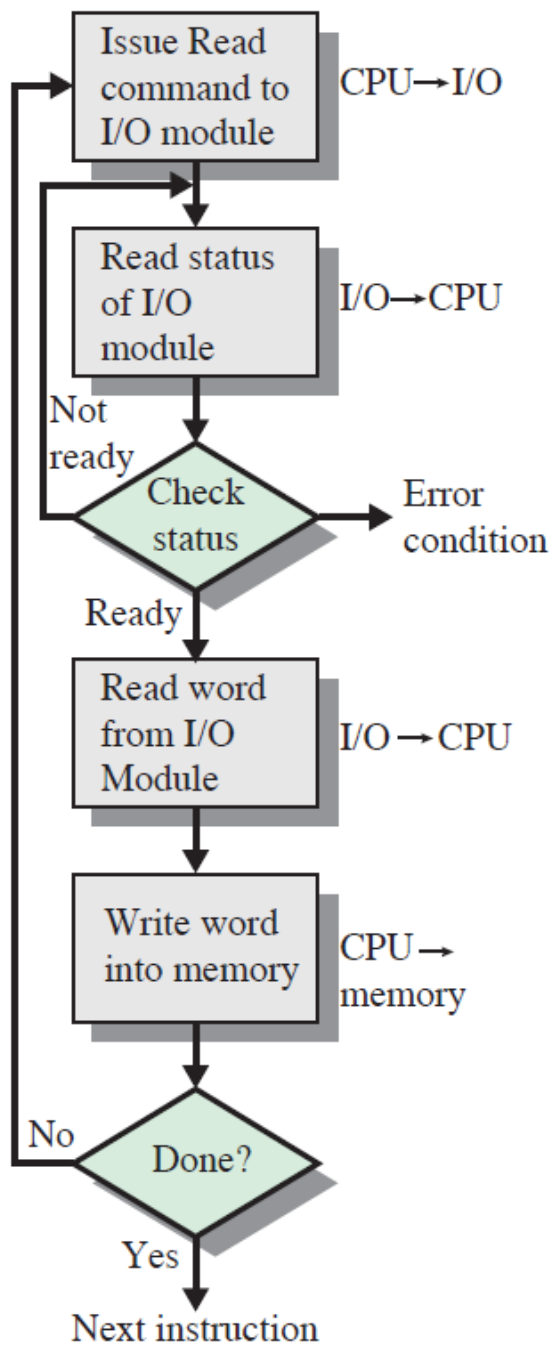
## ■ Interrupt-driven I/O

- the processor issues an I/O command on behalf of a process

## ■ Direct Memory Access (DMA)


- a DMA module controls the exchange of data **between main memory and an I/O module**






# Compare three techniques


- ❌ Transfer rate is **limited** by the speed with which the processor can test and service a device
- ❌ The processor is **tied up** in managing an I/O transfer, a number of instructions must be executed for **each** I/O transfer
- ✅ processor is involved only at the beginning and end of the transfer
- ✅ More efficient
- ❌ processor executes more slowly during a transfer when processor access to the **bus** is required




# Alternative DMA Configurations



(a) Single-bus, detached DMA



(b) Single-bus, Integrated DMA-I/O



(c) I/O bus

# Evolution of the I/O Function

1

- Processor **directly controls** a peripheral device

2

- A **controller** or **I/O module** is added

3

- Same configuration as step 2, but now **interrupts** are employed

4

- The I/O module is given direct control of memory via **DMA**

5

- The I/O module is enhanced to become a **separate processor**, with a specialized instruction set tailored for I/O

6

- The I/O module has a **local memory of its own** and is, in fact, a computer in its own right

# Design Objectives

---

## ■ Efficiency

- Major effort in I/O design
- Important because I/O operations often form a bottleneck
- Most I/O devices are extremely slow compared with main memory and the processor
- The area that has received the most attention is disk I/O

## ■ Generality

- Desirable to handle all devices **in a uniform manner**
- Applies to the way **processes view** I/O devices and the way the **operating system** manages I/O devices and operations
- **Diversity** of devices makes it difficult to achieve true generality
- Use a **hierarchical, modular approach** to the design of the I/O function

# Hierarchical Design


---

- **Functions of the OS should be separated according to their complexity, their characteristic time scale, and their level of abstraction**
- **Leads to an organization of the operating system into a series of layers**
- **Each layer performs a related subset of the functions required of the operating system**
- **Layers should be defined so that changes in one layer do not require changes in other layers**


# A Model of I/O Organization



(a) Local peripheral device



(b) Communications port



(c) File system

# Buffering

---

- **Perform input transfers in advance of requests being made and perform output transfers some time after the request is made**

## Block-oriented device

- stores information in blocks that are usually of fixed size
- transfers are made one block at a time
- possible to reference data by its block number
- disks and USB keys are examples


## Stream-oriented device

- transfers data in and out as a stream of bytes
- no block structure
- terminals, printers, communications ports, and most other devices that are not secondary storage are examples




# I/O Buffering Schemes (input)

---




(a) No buffering




(b) Single buffering

# I/O Buffering Schemes (input)



(c) Double buffering



(d) Circular buffering


# The Utility of Buffering

---

- Technique that **smoothes out peaks** in I/O demand
  - with enough demand eventually all buffers become **full** and their **advantage is lost**
- When there is a variety of I/O and process activities to service, buffering can
  - **increase the efficiency** of the OS
  - **increase the performance** of individual processes


# Diagram labeling the major components of HDD

---



# Disk Performance Parameters

- The actual details of disk I/O operation depend on
  - computer system
  - operating system
  - nature of the I/O channel and disk controller hardware



**Timing of a Disk I/O Transfer**

# Positioning the Read/Write Heads

---

- When the disk drive is operating, the disk is rotating at constant speed
- To read or write the head must be positioned at the desired track and at the beginning of the desired sector on that track
- Track selection involves moving the head in a movable-head system or electronically selecting one head on a fixed-head system
- On a movable-head system the time it takes to position the head at the track is known as **seek time**
- The time it takes for the beginning of the sector to reach the head is known as **rotational delay**
- The sum of the seek time and the rotational delay equals the **access time**

# An Example Sequence of I/O Requests

---

- we assume that


- the disk head is initially located at track 100
- a disk with 200 tracks
- the disk request queue has random requests in it

- the requested tracks, in the order received by the disk scheduler are

- 55, 58, 39, 18, 90, 160, 150, 38, 184

# First-In, First-Out (FIFO)

- Processes in sequential order
- Fair to all processes
- Approximates random scheduling in performance if there are many processes competing for the disk






# Priority (PRI)

---

- Control of the scheduling is outside the control of disk management software
- Goal is not to optimize disk utilization but to meet other objectives
- Short batch jobs and interactive jobs are given higher priority
- Provides good interactive response time
- Longer jobs may have to wait an excessively long time
- A poor policy for database systems


# Shortest Service Time First (SSTF)

- Select the disk I/O request that requires the least movement of the disk arm from its current position
- Always choose the minimum seek time




# SCAN

- Also known as the **elevator algorithm**
- Arm moves in **one direction only**
  - satisfies all outstanding requests until it reaches the last track in that direction then the direction is reversed
- Favors jobs whose requests are for tracks nearest to both innermost and outermost tracks



# C-SCAN (Circular SCAN)

- **Restricts** scanning to one direction only
- When the last track has been visited in one direction, the arm is returned to the **opposite end** of the disk and the scan begins again



# N-Step-SCAN

---





- Segments the disk request queue into subqueues of length  $N$
- Subqueues are processed one at a time, using SCAN
- While a queue is being processed new requests must be added to some other queue
- If fewer than  $N$  requests are available at the end of a scan, all of them are processed with the next scan

# FSCAN

---

- **Uses two subqueues**
- **When a scan begins, all of the requests are in one of the queues, with the other empty**
- **During scan, all new requests are put into the other queue**
- **Service of new requests is deferred until all of the old requests have been processed**

# Comparison of Disk Scheduling Algorithms



# Comparison of Disk Scheduling Algorithms

| (a) FIFO<br>(starting at track 100) |                                  | (b) SSTF<br>(starting at track 100) |                                  | (c) SCAN<br>(starting at track 100, in the<br>direction of increasing track<br>number) |                                  | (d) C-SCAN<br>(starting at track 100, in the<br>direction of increasing track<br>number) |                                  |
|-------------------------------------|----------------------------------|-------------------------------------|----------------------------------|--|----------------------------------|--|----------------------------------|
| Next track<br>accessed              | Number of<br>tracks<br>traversed | Next track<br>accessed              | Number of<br>tracks<br>traversed | Next track<br>accessed   | Number of<br>tracks<br>traversed | Next track<br>accessed   | Number of<br>tracks<br>traversed |
| 55                                  | 45                               | 90                                  | 10                               | 150  | 50                               | 150  | 50                               |
| 58                                  | 3                                | 58                                  | 32                               | 160  | 10                               | 160  | 10                               |
| 39                                  | 19                               | 55                                  | 3                                | 184  | 24                               | 184  | 24                               |
| 18                                  | 21                               | 39                                  | 16                               | 90   | 94                               | 18   | 166                              |
| 90                                  | 72                               | 38                                  | 1                                | 58   | 32                               | 38   | 20                               |
| 160                                 | 70                               | 18                                  | 20                               | 55   | 3                                | 39   | 1                                |
| 150                                 | 10                               | 150                                 | 132                              | 39   | 16                               | 55   | 16                               |
| 38                                  | 112                              | 160                                 | 10                               | 38   | 1                                | 58   | 3                                |
| 184                                 | 146                              | 184                                 | 24                               | 18   | 20                               | 90   | 32                               |
| Average seek<br>length              | 55.3                             | Average seek<br>length              | 27.5                             | Average seek<br>length   | 27.8                             | Average seek<br>length   | 35.8                             |



# Disk Scheduling Algorithms

| Name   | Description  | Remarks                                    |
|--|--|--|
| <b>Selection according to requestor</b>      |  |  |
| <b>RSS</b>                                   | Random scheduling  | For analysis and simulation                |
| <b>FIFO</b>                                  | First in first out   | Fairest of them all                        |
| <b>PRI</b>                                   | Priority by process  | Control outside of disk queue management   |
| <b>LIFO</b>                                  | Last in first out  | Maximize locality and resource utilization |
| <b>Selection according to requested item</b> |  |  |
| <b>SSTF</b>                                  | Shortest service time first                                | High utilization, small queues             |
| <b>SCAN</b>                                  | Back and forth over disk                                   | Better service distribution                |
| <b>C-SCAN</b>                                | One way with fast return                                   | Lower service variability                  |
| <b>N-step-SCAN</b>                           | SCAN of N records at a time                                | Service guarantee                          |
| <b>FSCAN</b>                                 | N-step-SCAN with N = queue size at beginning of SCAN cycle | Load sensitive                             |

# Summary

---

- I/O architecture is the computer system's interface to the outside world
- I/O functions are generally broken up into many layers
- A key aspect of I/O is the use of buffers
  - that are controlled by I/O utilities rather than by application processes
- Buffering smoothes out the differences between the speeds
- The use of buffers also decouples the actual I/O transfer from the address space of the application process
- Disk I/O has the greatest impact on overall system performance
- Two of the most widely used approaches are disk scheduling and the disk cache